

FINAL YEAR PROJECTAI ENGINEERINGCVPR CONFERENCE STYLE

ArSignTutor: Interactive Real-Time Arabic Sign Language Learning System

Using Deep Learning & Leakage-Aware Evaluation


STUDENT
Zeed Al Haj Ali

ID
2280205

FIELD
AI Engineering

DATE
June 2026

DATASET
RGB 7,856 Images



SECTION 01

WHY ARABIC SIGN LANGUAGE?

Community Need

7M+ Arabic-speaking deaf individuals with very few dedicated learning tools or interactive systems.

Research Gap

ASL dominates research — Arabic Sign Language remains severely underexplored despite community size.

Tool Deficit

Existing systems are static classifiers only. No interactive tutoring or real-time feedback for ArSL.

Evaluation Crisis

High benchmark scores hide leakage. Honest evaluation requires cross-dataset and cluster-based testing.

SECTION 04

PREPROCESSING & AUGMENTATION

RandomRotation
±15°

RandomAffine
Translate 8%, Scale 90-110%

ColorJitter
Brightness, Contrast, Saturation

RandomErasing
p=0.25, scale 2-12%

Normalization
ImageNet μ/σ (0.485/0.229...)

No Flip
Mirroring changes sign meaning

Applied to training only. Validation/test: resize + normalize only.

SECTION 06

TRAINING CONFIGURATION

Optimizer	AdamW
Loss	CrossEntropy + $\epsilon=0.10$
Batch Size	32
Epochs	Head: 5 Fine-tune: ≤25
LR	1e-3 → 1e-4
Scheduler	ReduceLROnPlateau
Early Stop	Patience = 6
Precision	AMP Mixed
Device	Kaggle Tesla T4
Init	ImageNet → ArASL → RGB

SECTION 02

DATASETS

ArASL Dataset

54,049 Images 32 Classes 40+ Contributors

⚠ Near-duplicate frames → random split causes data leakage

RGB Arabic Alphabets Dataset (Independent OOD)

7,856 Images 31 Classes 0 Leaky Clusters

✓ Clean cluster-based split: 5,509 train | 1,182 val | 1,165 test

SECTION 03

METHODOLOGY PIPELINE

Dataset Analysis dHash Cluster Clean Split Augment +Prep MBv2 Phase 1 Fine-tune Phase 2 Eval +Deploy

Phase 1 — Head Training		Phase 2 — Full Fine-Tuning	
Backbone	FROZEN X	Backbone	UNFROZEN ✓
Classifier	TRAINING ✓	Classifier	TRAINING ✓
LR	1e-3	LR	1e-4
Epochs	5	Epochs	≤25 + ES

SECTION 05

MODEL ARCHITECTURE — MobileNetV2

Input 224x224 Conv BN ReLU6 19xInverted Residual Conv BN ReLU6 Adaptive AvgPool FC → 31 Classes

1,280 Feature Dims 3.4M Parameters <50ms Inference 31 Classes

Init: ImageNet pretrained → ArASL fine-tuned → RGB domain-adapted

SECTION 12

MAIN CONTRIBUTIONS

Developed ArSignTutor educational interactive platform for Arabic Sign Language

Applied MobileNetV2 transfer learning with two-phase fine-tuning strategy

Introduced leakage-aware dHash cluster-based evaluation protocol

Performed cross-dataset validation exposing 99%→48.78% generalization gap

Achieved 94.33% clean test accuracy — scientifically defensible result

SECTION 07

EXPERIMENTAL RESULTS


TEST ACCURACY 94.33% PRECISION 94.59% RECALL 94.33% F1-SCORE 94.36%

SECTION 08

GENERALIZATION ANALYSIS

Original Model	Random (Image)	ArASL	99.26% LEAKY
Sanity Check	—	ArASL	98.75% PIPELINE OK
Cross-Dataset OOD	ArASL → RGB	RGB 31cls	48.78% GAP
Clean Retrain ✓	Cluster (dHash)	Held-out RGB	94.33% HONEST


48.78% OOD drop = leakage + domain shift + class mapping. Clean 94.33% is the honest result.



SECTION 13

DEPLOYMENT

Framework	Streamlit
Model	MobileNetV2
Backend	PyTorch
Inference	Real-Time <50ms
Platform	Desktop Prototype
Classes	31 Arabic Letters



SECTION 14

SYSTEM DEMONSTRATION

Main Interface — Spelling Tutor

Letter Recognition Mode — Real-Time Webcam

SECTION 11

LIMITATIONS & FUTURE WORK

Static Signs Only

Dynamic word-level signs need video/temporal models

31 of 32 Letters

yaa excluded — no RGB samples available

Single-Frame

Temporal smoothing only partially fixes jitter

SECTION 15

CONCLUSION

ArSignTutor demonstrates that Arabic Sign Language learning can be effectively supported through an interactive deep learning system.

After removing data leakage via dHash cluster-based evaluation and domain adaptation to RGB data, the final model achieved 94.33% clean test accuracy — a scientifically honest and reproducible result.

The key lesson: a near-perfect benchmark score is not always a good model. Checking the evaluation protocol matters as much as the architecture.

- Video temporal models for word-level signs
- MediaPipe crop retraining
- Grad-CAM visualization
- Mobile deployment (Android/iOS)